

Algorithms, Design & Analysis

Lecture 22: Longest Path Problem

Afrazia Umer Farooq & Qudsia Khan

Information Technology University

April 22, 2025

About Your Fellows

- Hello guys! **Afrazia** and **Qudsia Khan** here.
- We are Associate Students at ITU.

Overview

- Recap: Shortest Path Algorithms
- Longest Path Problem Statement
- NP-Hardness and Complexity
- Special Case: DAGs Solution
- Example Walkthrough

Recap: Shortest Path Algorithms

- **Dijkstra's:** Non-negative weights
- **Bellman-Ford:** Handles negative weights
- **Floyd-Warshall:** All-pairs shortest paths
- **DFS:** For unweighted graphs

Longest Path Problem

(Finding the Longest Path in a Directed Acyclic Graph (DAG))

Longest Path Problem Statement

Definition

Find a path in graph $G = (V, E, W)$ with maximum sum of edge weights

- Weights are non-negative
- Works for both directed/undirected graphs
- Start and end nodes can be arbitrary
- Only simple paths (no repeated nodes)

Characteristics

Key Characteristics:

- Path: A sequence of vertices with directed edges between them.
- Longest Path: Path with maximum total weight.
- DAGs are efficiently solvable due to absence of cycles.
- No vertex will be used more than once i.e no repeated vertex.

Difference between normal graphs and DAGs:

- **For general graphs:** Longest path problem is NP-hard because it does not have optimal structure property.
- Time complexity in general graphs is \mathbf{NP}
 - **NP** (Non-deterministic Polynomial time): Class of problems where a solution can be **verified** in polynomial time.
 - **NP-Hard** problems are **at least as hard as the hardest NP problems**.
 - **No known fast (polynomial-time) algorithm** exists to solve NP-Hard problems in all cases.
- **For directed acyclic graphs (DAG):** Longest path problem has **linear time complexity**.

Solving Strategies

Solving Strategies:

There are two main approaches to solve the Longest Path Problem:

- **Topological Sorting Algorithm**
- **Dynamic Programming**

Topological Sorting Algorithm


Topological Sorting Algorithm

- **Topological sort:**

- Since DAG has no cycles, topological sort is possible.
- Topological order ensures that for every edge $(u \rightarrow v)$, u appears before v .

- **Example:**

Let us consider the following graph:



Topological Sorting Algorithm

Distance from Source Node 'c'

Node	Distance
a	$-\infty$
b	$-\infty$
c	0
d	$-\infty$
e	$-\infty$
f	$-\infty$

- Source node **c** initialized to 0
- All other nodes set to $-\infty$ (unreachable)
- Will update distances during relaxation

Step 1: Start at Node c

Cumulative Weight at Node c:

- From node **c**, there is an edge to node **b** with weight 4.
- Cumulative weight: 4

Updated DP table:

$$\text{dp}[c] = 0, \quad \text{dp}[b] = 4, \quad \text{dp}[a] = -\infty, \quad \text{dp}[e] = -\infty, \quad \text{dp}[d] = -\infty, \quad \text{dp}[f] = -\infty$$

Step 2: Process Node b

Cumulative Weight at Node b:

- From node **b**, there is an edge to node **a** with weight 1.
- Cumulative weight: $4 + 1 = 5$

Updated DP table:

$$\text{dp}[c] = 0, \quad \text{dp}[b] = 4, \quad \text{dp}[a] = 5, \quad \text{dp}[e] = -\infty, \quad \text{dp}[d] = -\infty, \quad \text{dp}[f] = -\infty$$

Step 3: Process Node a

Cumulative Weight at Node a:

- From node **a**, there is an edge to node **e** with weight 2.
- Cumulative weight: $4 + 1 + 2 = 7$

Updated DP table:

$$\text{dp}[c] = 0, \quad \text{dp}[b] = 4, \quad \text{dp}[a] = 5, \quad \text{dp}[e] = 7, \quad \text{dp}[d] = -\infty, \quad \text{dp}[f] = -\infty$$

Step 4: Process Node e

Cumulative Weight at Node e:

- From node **e**, there is an edge to node **d** with weight 5.
- Cumulative weight: $4 + 1 + 2 + 5 = 12$

Updated DP table:

$$dp[c] = 0, \quad dp[b] = 4, \quad dp[a] = 5, \quad dp[e] = 7, \quad dp[d] = 12, \quad dp[f] = -\infty$$

Step 5: Process Node d

Cumulative Weight at Node d:

- From node **d**, there is an edge to node **f** with weight 3.
- Cumulative weight: $4 + 1 + 2 + 5 + 3 = 15$

Updated DP table:

$$dp[c] = 0, \quad dp[b] = 4, \quad dp[a] = 5, \quad dp[e] = 7, \quad dp[d] = 12, \quad dp[f] = 15$$

Step 6: Process Node f

Cumulative Weight at Node f:

- Node **f** has no outgoing edges.
- Cumulative weight: 15 (final weight)

Final DP table:

$$\text{dp}[c] = 0, \quad \text{dp}[b] = 4, \quad \text{dp}[a] = 5, \quad \text{dp}[e] = 7, \quad \text{dp}[d] = 12, \quad \text{dp}[f] = 15$$

Longest Path

Path Weights :

- $c \rightarrow b = 4$ (4)
- $b \rightarrow a = 5$ (4+1)
- $a \rightarrow e = 7$ (4+1+2)
- $e \rightarrow d = 12$ (4+1+2+5)
- $d \rightarrow f = 15$ (4+1+2+5+3)

Longest Path:

Total: 15 in $O(|V| + |E|)$
 $c \rightarrow b \rightarrow a \rightarrow e \rightarrow d \rightarrow f$

- Path shown with cumulative weights at each step
- Final path length matches your calculation
- **Time complexity** is $O(|V| + |E|)$

Dynamic Programing

Longest Path in DAG —DP (Step 1)

- **Goal:** Find the longest path in a Directed Acyclic Graph (DAG).
- **Step 1: Initialize**
 - For each node, create $lp[node] = 0$.
 - If a node has no outgoing edges, its longest path remains 0
- **Step 2: Maintain Incoming Edge Count**
 - Count how many incoming edges each node has.
 - This helps us know when a node is ready to be processed.
- **Step 3: Start Queue**
 - Add all nodes with no outgoing edges into the processing queue.

Longest Path in DAG — DP (Step 2)

• Step 4: Process Nodes from Queue

- While the queue is not empty:
- Take a node u from the queue.
- For every node v that goes into u (i.e., there is an edge $v \rightarrow u$):
 - Update: $lp[v] = \max(lp[v], lp[u] + \text{weight}(v,u))$
 - Decrease incoming edge count of v .
 - If v now has no more outgoing edges, add it to the queue.


• Step 5: Final Answer

- After the queue is empty, the longest path in the DAG is the maximum value in $lp[]$ array.

• Time Complexity: $\mathcal{O}(V + E)$

Problem Overview

- **Goal:** Find the longest path in a Directed Acyclic Graph (DAG).



Initialization (Destination Node)

Goal: Find longest path that ends at node f

Initialization:

- $lp[f] = 0$ (start from destination)
- All other lp values $= -\infty$
- Initial queue $= [f]$

Processing Node f

Queue before: $[f]$

Pop: f

Incoming edges:

- $d \rightarrow f$ (weight 3)
- $c \rightarrow f$ (weight 2)

Updates:

- $lp[d] = 0 + 3 = 3$
- $lp[c] = 0 + 2 = 2$

Queue after: $[d, c]$

Processing Node d

Queue before: [d, c]

Pop: d

Incoming edges:

- $e \rightarrow d$ (weight 5)
- $c \rightarrow d$ (weight 3)

Updates:

- $lp[e] = 3 + 5 = 8$
- $lp[c] = \max(2, 3 + 3) = 6$

Queue after: [c, e]

Processing Node c

Queue before: [c, e]

Pop: c

Incoming edges: None

No updates

Queue after: [e]

Processing Node e

Queue before: [e]

Pop: e

Incoming edges:

- $a \rightarrow e$ (weight 2)
- $b \rightarrow e$ (weight 3)

Updates:

- $lp[a] = 8 + 2 = 10$
- $lp[b] = 8 + 3 = 11$

Queue after: [a, b]

Processing Node a

Queue before: [a, b]

Pop: a

Incoming edge: $b \rightarrow a$ (weight 1)

Update: $lp[b] = \max(11, 10 + 1) = 11$ (no change)

Queue after: [b]

Processing Node b

Queue before: [b]

Pop: b

Incoming edge: $c \rightarrow b$ (weight 4)

Update: $lp[c] = \max(6, 11 + 4) = 15$

Queue after: [c]

Processing Node c (again)

Queue before: [c]

Pop: c

Incoming edges: None

No updates

Queue after: [] (empty)

Final lp Values

- $lp[a] = 10$
- $lp[b] = 11$
- $lp[c] = \mathbf{15}$
- $lp[d] = 3$
- $lp[e] = 8$
- $lp[f] = 0$

Final Answer

- Longest path ending at f: **15**
- Path: $c \rightarrow b \rightarrow a \rightarrow e \rightarrow d \rightarrow f$

Final Answer

- Longest path ending at f: **15**
- Path: $c \rightarrow b \rightarrow a \rightarrow e \rightarrow d \rightarrow f$

Time Complexity

- **Time Complexity of the Longest Path Algorithm (DP):**

- **Reverse Traversal:** Each node's incoming edges are checked once, so total edge processing is $O(E)$.
- **Distance Updates:** Each node is processed once and updated based on its incoming neighbors.
- **Queue Operations:** Each node and edge can contribute at most once to the queue, maintaining $O(V + E)$ complexity.
- **Overall Time Complexity:** The dynamic programming approach still maintains:

$$O(V + E)$$

Conclusion

- **Recap:**

- We used Dynamic Programming starting from node f and propagated backward through incoming edges.
- The longest path ending at f is 15, and the correct path is $c \rightarrow b \rightarrow a \rightarrow e \rightarrow d \rightarrow f$.

- **Key Takeaways:**

- DP can be used when the endpoint (like f) is known and incoming edges are easy to trace.
- The queue simulates the flow of longest distance calculations backward through the graph.
- The time complexity remains $O(V + E)$, making this approach efficient for DAGs.

Relationship between Dynamic Programming questions and DAG

Why is any DP problem a DAG?

- Each subproblem depends only on smaller subproblems.
- The dependencies move in one direction — from smaller to larger.
- We can draw this as a graph:
 - Nodes = subproblems
 - Edges = dependencies
- There are no cycles — no subproblem depends on itself.

Therefore:

The structure of subproblems and their dependencies forms a **Directed Acyclic Graph (DAG)**.

Fibonacci as a DAG

Fibonacci recurrence:

$$F(n) = F(n - 1) + F(n - 2)$$

- To compute $F(5)$, we must compute $F(4)$ and $F(3)$.
- To compute $F(4)$, we need $F(3)$ and $F(2)$.
- To compute $F(3)$, we need $F(2)$ and $F(1)$.
- This dependency continues downward until base cases.

Insight

These recursive calls form a DAG — every node depends on smaller nodes, and there are no cycles.


Problem Statement

- Compute $F(5)$ using reverse dynamic programming (bottom-up).
- Fibonacci is defined as:

$$F(0) = 0, \quad F(1) = 1, \quad F(n) = F(n-1) + F(n-2) \text{ for } n \geq 2$$

- We'll process in order: $F(2), F(3), F(4), F(5)$, tracking computation and queue at each step.

Initialization



Initialization:

$$F(0) = 0, \quad F(1) = 1$$

Queue: [2, 3, 4, 5]

Processing Node: $F(2)$


Queue before: [2, 3, 4, 5]

Pop: $F(2)$

Dependencies: $F(1) = 1, F(0) = 0$

$$F(2) = F(1) + F(0) = 1 + 0 = 1$$

Queue after: [3, 4, 5]



Processing Node: $F(3)$


Queue before: [3, 4, 5]

Pop: $F(3)$

Dependencies: $F(2) = 1, F(1) = 1$

$$F(3) = F(2) + F(1) = 1 + 1 = 2$$

Queue after: [4, 5]



Processing Node: $F(4)$


Queue before: $[4, 5]$

Pop: $F(4)$

Dependencies: $F(3) = 2, F(2) = 1$

$$F(4) = F(3) + F(2) = 2 + 1 = 3$$

Queue after: $[5]$



Processing Node: F(5)


Queue before: [5]

Pop: F(5)

Dependencies: $F(4) = 3$, $F(3) = 2$

$$F(5) = F(4) + F(3) = 3 + 2 = \boxed{5}$$

Queue after: []



Final Fibonacci Table

F(i)	Value
F(0)	0
F(1)	1
F(2)	1
F(3)	2
F(4)	3
F(5)	5

Conclusion

- Each Fibonacci value was treated as a node in a DAG.
- The dependencies formed edges, and the problem was solved in bottom-up order using DP.
- This method avoids recomputation and mirrors DAG traversal with queue.
- Final result:

$$F(5) = 5$$

Practice Questions

Rod Cutting Problem

Problem:

Given a rod of length n inches and an array of prices that contains prices of all pieces of size smaller than n , determine the maximum value obtainable by cutting up the rod and selling the pieces.

Input: Integer n , and array $price[]$ of size n

Output: Maximum total price from cutting and selling the rod

Recurrence Relation:

$$DP[i] = \max_{1 \leq j \leq i} (price[j - 1] + DP[i - j])$$

Painting Fence Problem

Problem:

Given n posts and k colors, find the number of ways to paint the posts such that no more than two adjacent posts have the same color.

Input: Integers n and k

Output: Number of valid ways to paint

Recurrence Relation:

$$DP[i] = (k - 1) \times (DP[i - 1] + DP[i - 2])$$


Max Coin Collection Grid Game

Game Summary

- You are given a grid with hidden coins placed on certain cells.
- There are two extractors:
 - Extractor — starts at bottom-left and moves **up**.
 - Extractor — starts at bottom-left and moves **right**.
- Allowed moves: move either extractor one step in its direction.
- When an extractor moves to a new row or column containing coins:
 - One coin is collected.
 - All other coins in that row or column are lost.
- The game ends when either extractor reaches the top or right edge.
- **Goal:** Plan your extractor moves to collect the **maximum number of coins**.

Initial Coin Grid


Start at (0,0). Coins (●), blocked cells (X):



No coins collected yet.

First Move


Pick coin at (1,3). Block row 1 and column 3.



Coins collected: (1,3)

Second Move


Pick coin at (3,4). Block row 3 and column 4.



Coins collected: (1,3), (3,4)

Third Move

Pick coin at (4,5). Block row 4 and column 5.



Coins collected: (1,3), (3,4), (4,5)

Game Summary

- Total coins collected: **3**
- Other coins blocked due to row/column restrictions.
- You reached a dead-end – further coins not accessible.

Q: When does a greedy approach fail in this game?